

# JEEVA VINCENT

## AI-Native Data Engineer

+91 99621 26629 | jeevavincent.2003@gmail.com | Thoothukudi, Tamil Nadu, India

LinkedIn | GitHub



### PROFESSIONAL SUMMARY

Full Stack AI Engineer at iCustomer with over a year of progressive experience — promoted twice from Junior Data Engineer to AI-Native Data Engineer to Full Stack AI Engineer. Specializes in building end-to-end AI-powered data systems, LLM-based enrichment services, agentic scoring engines, full-stack platforms, and ML-driven segmentation systems across GCP and AWS. Passionate about data quality, automation, and leveraging LLMs to solve real-world B2B data challenges.

### TECHNICAL SKILLS

|                       |                                                                                                                                                                      |
|-----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Languages             | Python, Go, SQL, TypeScript, JavaScript                                                                                                                              |
| Backend Frameworks    | FastAPI, NestJS, Express.js, Cobra (Go CLI), Mangum (AWS Lambda adapter)                                                                                             |
| Frontend & UI         | React, Next.js, Streamlit, Tailwind CSS, Framer Motion                                                                                                               |
| Data Engineering      | Google BigQuery, PostgreSQL, Snowflake, dbt, Apache Airflow, Apache Kafka, RabbitMQ, Apache Spark, Meltano (ELT), dlt, Airbyte, Pyarrow                              |
| AI / ML               | Anthropic Claude (Sonnet, Haiku), OpenAI GPT-4, Pydantic AI, K-Means Clustering, RFM Modelling, XGBoost, scikit-learn, sentence-transformers, LLM Prompt Engineering |
| Cloud & DevOps        | GCP (BigQuery, GCS, Cloud Run), AWS (Lambda, S3, EC2, Glue, Athena, ECR), Docker, Boto3, Redis                                                                       |
| Databases & ORM       | PostgreSQL, BigQuery, Snowflake, SQLAlchemy, TypeORM, asyncpg, psycopg2, Qdrant (Vector DB)                                                                          |
| Data Enrichment       | Zyte API, BrightData (SERP), BuiltWith / Wappalyzer, ReverseContact, Pubrio, Apollo.io, Lusha, Dropcontact, ZeroBounce                                               |
| Scraping & Automation | Selenium, BeautifulSoup, aiohttp, httpx, Playwright, Requests, lxml                                                                                                  |
| Platforms & Analytics | OneSource (iCustomer Platform), LookML / Looker, Salesforce (SFDC), Slack SDK, Plotly, Matplotlib, Seaborn, Pandas, NumPy                                            |
| Security & Auth       | AES-256-CBC Encryption (PyCryptodome), JWT, OAuth 2.0, AWS Secrets Manager                                                                                           |

### WORK EXPERIENCE

#### Full Stack AI Engineer

July 2026 – Present

iCustomer — Cambridge, MA (Hybrid)

- Building and owning full-stack AI-powered features — from data pipeline backends (Python/FastAPI) to frontend dashboards (React/Next.js) and AI agent integrations.
- Designing and deploying agentic systems using Anthropic Claude and OpenAI GPT-4 for intelligent B2B data processing, scoring, and enrichment at scale.
- Leading the iConnector platform — a universal data sync and activation service built as both a Python (PyPI) and npm package, integrating 43 connectors including HubSpot, Salesforce, Klaviyo, LinkedIn Ads, and Meta Ads using the Airbyte CDK.
- Architecting multi-tenant SaaS data infrastructure on GCP and AWS — Cloud Run, BigQuery, PostgreSQL, Redis, BullMQ job queues.

#### AI-Native Data Engineer

February 2026 – June 2026

iCustomer — Cambridge, MA (Hybrid)

- Designed and deployed the FIRE Scoring Engine (Fit, Intent, Recency, Engagement) in Python — triggered via RabbitMQ, orchestrated dbt transformations on BigQuery, persisted account-level scores to PostgreSQL with a configurable Streamlit UI.
- Built a two-agent ICP Scoring Engine using Anthropic Claude Sonnet + Haiku: Agent 1 parses unstructured ICP documents (PPTX, DOCX, PDF, XLSX) into structured templates; Agent 2 scores accounts in parallel batches outputting fit scores (0–100) with segment classification.
- Delivered end-to-end RFM K-means Customer Segmentation for Todd Snyder / American Eagle Outfitters — engineered models on 83K+ Shopify orders, froze trained model into a FastAPI prediction service, automated daily rescoring via Apache Airflow DAGs.
- Engineered DJ Graffiti — a fully automated weekly event discovery pipeline scraping 1,000s of company websites, classifying upcoming events using LLM extraction, delivering structured records into the data warehouse every Monday.

- Architected a Provider Orchestrator — intelligent enrichment routing with cascading fallback logic across ReverseContact and Pubrio for zero-downtime provider swaps.
- Built a Unified Enrichment Agent integrating OpenAI GPT-4, Zyte API, and BrightData to automatically enrich companies with identity, financial, technographic, and compliance data.
- Developed a Snowflake Native App (iCustomer CDO) with Streamlit UI and Snowpark Python ETL for deterministic identity resolution and B2B audience building inside customer Snowflake environments.
- Built the OneSource Platform backend in Go — chi HTTP router, multi-provider waterfall enrichment engine, 20+ CLI commands covering enrichment, identity resolution, lead generation, CRM sync.

## Junior Data Engineer

June 2025 – January 2026

iCustomer — Cambridge, MA (Hybrid)

- Developed a B2B Realtime Signal Tracker — async FastAPI service using concurrent aiohttp to scrape company websites, detect tech stacks, extract social handles and job openings, and rank accounts by engagement priority in real time.
- Built a Tag Identification Microservice deployed on AWS Lambda — detects tech stacks from a fingerprint database of 1,000+ app signatures, extracts social handles, job opening signals, and company descriptions from live websites.
- Implemented the OneSource Contact Discovery API on AWS Lambda: queries PostgreSQL and BigQuery, applies AES-256-CBC decryption for secure key handling, and returns enriched contact profiles via REST.
- Built an RB2B Webhook receiver for real-time anonymous B2B visitor de-anonymization — processes incoming payloads and inserts structured visitor records into per-tenant BigQuery tables instantly.
- Ran large-scale domain validation and technographic enrichment pipelines — split 50K+ domain lists into parallel batches and enriched tech stack data via BuiltWith integration.
- Authored LookML analytics views in Looker covering monthly customer analytics, cohort retention, LTV over time, company benchmarks, and cluster summaries.
- Integrated Meltano ELT platform with a custom React / TypeScript management UI for configuring extractors, loaders, and schedules with live pipeline monitoring.
- Built an SEO Keywords Finder (Async FastAPI on AWS Lambda) — extracts website meta content and generates B2B enrichment keywords using OpenAI GPT.

## PERSONAL PROJECTS

### Tech Intelligence Pipeline

Python, Kafka, Apache Airflow, dbt, DuckDB, FastAPI

End-to-end market intelligence pipeline ingesting live crypto/stock data, Reddit posts, and news articles via Kafka producers, transforming with dbt, storing in DuckDB, and exposing sentiment and trend APIs.

### Sentinel Shield — Personal Network Privacy Guard

Python, Scapy, Bytewax, Ollama (local LLM), LanceDB, FastAPI

Real-time network intrusion detection system that captures packets with Scapy, processes the stream with Bytewax, analyzes traffic with a local AI model, stores embeddings in LanceDB, and sends instant notifications.

### Real-Time Stock Market Data Pipeline

Apache Kafka, Python, AWS (S3, Glue, Athena)

Kafka-based ETL pipeline streaming stock market data from CSV to JSON, storing in AWS S3, cataloging with AWS Glue, and querying trends via Athena.

### Uber Data Engineering Pipeline

Python, Apache Spark, Apache Airflow, Google BigQuery

Batch processing pipeline analyzing Uber ride data with Spark, orchestrated by Airflow DAGs, generating insights on peak hours, location heatmaps, and revenue trends in BigQuery.

## EDUCATION

### Master of Business Administration (MBA)

2025 – 2027 (Pursuing)

Bharathidasan University, Tiruchirapalli, Tamil Nadu (Distance Education)

Specialization: Human Resource Management

### Bachelor of Computer Applications (BCA)

August 2022 – May 2025

Providence College for Women (Autonomous), Coonoor, Tamil Nadu

GPA: 8.47 / 10.00 | Bharathiar University Affiliation

### Diploma in Cyber Security (DCS)

August 2022 – May 2025

Bharathiar University, Coimbatore, Tamil Nadu